

Educational Measurement as an AI-Native Profession

Session format: Panel presentation, 90 minutes (4 panelists and a discussant; AI-moderated under a human session chair)

Session: Wednesday, 7 October 2026, 11:00 AM–12:30 PM, Commonwealth 2

Conference: AIME-Con 2026, 5–7 October 2026, Wyndham Grand Pittsburgh Downtown

Topic of interest: AI-integrated professional practice in educational measurement

Document status: Final, 22 September 2026

This is the final session proposal. The Abstract, Description, Format, and Structure are the accepted text, revised only where the change of discussant and the reordering of the guiding questions required it. The Addendum describes how the AI interlocutor, Iain, will take part, and says plainly which parts of that design are in place and which are still being built.

Contributors

Session chair Damian Betebenner, Center for Assessment

The chair holds final authority over every AI intervention and retains override and stop authority throughout the session.

Moderator Iain, an AI interlocutor operating under the session chair's supervision

Moderation is performed by an AI under human supervision; accountability remains human. The name is set lower case in the program and everywhere else. See the Addendum.

Panelists

- Damian Betebenner, Center for Assessment
- Derek Briggs, University of Colorado Boulder
- Frank Rijmen, Cambium Assessment
- Mohammed A. A. Abulela, MetaMetrics, Inc. / University of Minnesota
with co-authors: Guher Gorgun, University of Georgia; Brian French, Washington State University; Brian Leventhal, James Madison University; Matthew Gushta, MetaMetrics, Inc.

Discussant Fred Oswald, University of California, Irvine

Artificial intelligence is not only changing assessment products; it is changing the professional work of educational measurement itself. This panel examines how an AI-native future may transform day-to-day practice, professional norms, and epistemic responsibility in the field. Bringing together perspectives from measurement consulting, higher education, assessment industry, and graduate training, the panel will address AI-integrated psychometric and data-science workflows, scholarly publishing and peer review, assessment development and operations, graduate training and curricula, documentation, and governance. The session emphasizes concrete use cases rather than general speculation: how AI is already changing technical work, where it can responsibly accelerate, critique, or audit professional judgment, and where human accountability must remain primary. The discussion will focus on emerging standards for validity, fairness, transparency, reproducibility, attribution, quality control, and public trust as AI becomes embedded in measurement practice. The goal is to move from tool-centered conversations toward a profession-level account of how educational measurement should adapt while preserving rigor and human responsibility.

Session Description

Description + Format + Structure: 1000 / 1000 words

Artificial intelligence is becoming part of the infrastructure through which educational measurement work is conducted. The field has begun to examine AI for item generation, automated scoring, formative feedback, adaptive assessment, and classroom support. These are important, but they do not exhaust the professional implications of AI. A more fundamental question: how will AI transform the work, norms, and responsibilities of measurement professionals themselves?

This panel asks whether educational measurement needs a discipline-level response to an AI-native future, as the mathematics community recently has.¹ Wherever AI enters psychometric modeling, documentation, code, assessment design, review, and policy analysis, the field faces the questions other technical communities are already asking about the values that define them: standards of evidence, peer review, authorship, attribution, reproducibility, and public trust.

The panel therefore focuses less on AI as a product feature than as a transformation in professional practice: on how it will alter (and is already altering) the way measurement professionals reason, write, program, review, validate, and decide.

Fred Oswald will serve as discussant. Professor of Education at the University of California, Irvine, editor-in-chief of *Psychological Methods*, past president of the Society for Industrial and Organizational Psychology, past chair of the APA Committee on Psychological Tests and Assessment and of the National Academies Board on Human-Systems Integration—which

¹See the *Leiden Declaration on Artificial Intelligence and Mathematics* (2 June 2026; DOI: 10.5281/zenodo.20302944), a community initiative endorsed by the International Mathematical Union, which calls on mathematicians to protect the discipline's core values—the independent verifiability of proof, attribution, authorship, and the primacy of human judgment—as AI becomes embedded in mathematical research, publishing, and peer review: <https://leidendeclaration.ai/>.

produced *Human-AI Teaming*—and a former member of the National Artificial Intelligence Advisory Committee. He will synthesize the provocations and, based on them, press toward actionable norms, while critiquing the session's own AI moderation process. He will also press the reciprocal question of whether psychometric standards of validity, reliability, and fairness should govern how the field evaluates AI systems.

Damian Betebenner will discuss how AI has transformed his work as a measurement methodologist and data scientist, contrasting pre-AI and AI-integrated workflows across theoretical development, coding, simulation, data analysis, documentation, and the design of systems connecting achievement, growth, and improvement. The emphasis is on AI as a cognitive collaborator that changes the feasible scale, speed, and ambition of the work while creating new demands for verification and professional judgment.

Derek Briggs will focus on scholarly publishing, peer review, and graduate training. He will examine whether AI should become part of the manuscript-review infrastructure by generating systematic first-pass reviews keyed to journal scope, methodological standards, and reporting expectations, while human reviewers audit, correct, and extend those critiques. He will also address norms for author disclosure and for teaching graduate students to use AI without outsourcing judgment, originality, or methodological responsibility.

Frank Rijmen will bring an assessment-industry perspective to a question that has drawn less attention than AI's effects on specific processes and products: how AI is reshaping the archetypal roles of the psychometrician. He will frame the argument around three professional archetypes. The *Guardian* treats the *Standards* as first principles and is skeptical of any innovation that cannot be justified within the established framework. The *Catcher* is defined by operational vigilance. Meticulous and loss-averse, the *Catcher* focuses on what must not go wrong in calibration, equating, scoring, and anomaly detection. The *Architect* bridges psychometrics with data science and emerging technology, restless to redesign rather than execute.

Rijmen will argue that AI changes not only the knowledge and skills these roles require but also their non-cognitive demands: the dispositions and traits that make someone good at the work. He will consider what this means for how organizations define psychometric expertise, hire and develop staff, and prepare the next generation of measurement professionals.

Mohammed A. A. Abulela will address integrating AI into educational measurement curricula. AI is rapidly reshaping educational measurement, psychometrics, assessment development, data analytics, and the job market for measurement specialists and psychometricians, yet little is known about how graduate programs have responded. Programs should prepare future measurement professionals to use, evaluate, and responsibly integrate AI in nearly every aspect of their work. His contribution draws on a review of coursework across 90 North American doctoral programs in educational measurement, quantitative methodology, statistics, and related fields, plus an initial review of academic and industry job postings indicating growing demand for AI-related skills. Very few programs offered qualifying AI- or computationally related coursework, with such coursework identified in only 16 of 90 programs (17.8%). In the 10 programs where curricular status could be determined, elective offerings outnumbered required ones (6 versus 4), and all four programs with required coursework were explicitly oriented toward statistics and/or data science. He will discuss curricular gaps, workforce expectations, responsible AI use, and practical strategies for integrating AI-related knowledge and skills into graduate curricula.

Session Format

The session is an interactive, AI-moderated panel rather than a sequence of papers, enacting the practices it examines. The panelists' materials—slides, papers, and code—form a shared corpus beforehand. During the session Iain, an AI interlocutor supervised by the chair and grounded in that corpus, maintains a running record of claims, evidence, agreements, and open tensions; surfaces source-grounded questions; and distinguishes what it has observed, retrieved, and inferred. Its authority to speak is bounded to human-approved moments; humans retain final judgment. Live transcription of the session is in place; the rest of Iain's role will be used only where it proves reliable in rehearsal (Addendum, §1).

Five guiding questions organize the discussion, and Iain uses them to track coverage:

1. **What has changed.** How has AI already changed educational measurement work—in the substance, scope, and organization of that work, not only its speed—and what has become better, worse, or simply different?
2. **What that demands of professionals.** As AI becomes a collaborator rather than a tool, how should work be divided between people and AI; what knowledge, skills, and dispositions then define psychometric expertise; and how should graduate programs and employers develop and assess them?
3. **What the profession must redefine.** When AI contributes materially to an analysis, a manuscript, or a review, what must be verified, documented, and disclosed before others can responsibly rely on it—and what does that ask of peer review, training, and institutional accountability?
4. **What measurement science owes the evaluation of AI.** What can measurement science contribute, in return, to the evaluation of AI systems and of human-AI teams, including this panel, and what would count as improvement over the relevant alternative?
5. **Where authority and accountability remain human.** Where should authority and accountability remain human as capability advances, what justifies those boundaries beyond current capability, and what does this session's own use of an AI interlocutor reveal about them?

The last question comes last because it is the hardest, and the field's usual answers are not answers. Asserting that humans must retain final say does not yet supply the justification for it; observing that AI cannot yet do something states a fact with a short shelf life. The panel will press instead for grounds that survive advances in capability: who holds the warrant to make a claim, who is accountable to the people a decision affects, and who answers when it is wrong. Competent performance does not by itself confer authority, and retaining accountability does not require producing every part of the work by hand. Because Iain takes part under explicit limits, the session supplies a case rather than a hypothesis: not only what the AI could do, but what it was permitted to do, what it should have done, and when it rightly abstained.

The questions are a structure for the discussion, not five rounds. Question 1 is carried by the opening provocations; questions 2 through 5 organize the exchanges that follow; and

the norm-building exercise draws mainly on 3 and 5. Each candidate norm should name a concrete activity, the conditions under which AI involvement is acceptable, and who is answerable for the resulting decision.

Session Structure

90 minutes: 6 minutes for framing and an explanation of the AI moderation and its governance; 20 minutes for four five-minute provocations, each making one claim, one example, and one unresolved problem; 5 minutes for an AI synthesis of convergence, disagreement, and cross-panel connections; 12 minutes for Oswald's response, critiquing the panelists and the AI synthesis; 33 minutes of AI-moderated panel and audience dialogue; 11 minutes of norm-building in which Iain proposes candidate norms that panelists accept, revise, or reject; and 3 minutes for closing synthesis and a durable session record.

The session extends the conference theme, "Measurement Science in AI-Integrated Assessment and Pedagogy," from AI-integrated assessment systems to AI-integrated measurement work: to identify norms that let the field use AI to improve the quality, scope, and speed of that work while preserving evidentiary standards, human judgment, and public trust.

Session Design: How Iain Participates

1. What is in place, and what is still being built

This addendum describes the full design. Not all of it will be ready in October, and we would rather say so now than let the room discover it. The session makes three levels of promise:

Level	What it covers
Committed	A human-chaired panel built on the five guiding questions and the norm-building exercise. Live speech-to-text transcription of the session: the house audio system will give us a line feed into the capture machine, and a return path back into the room. A timekeeping display.
Planned	Iain's running record on screen (§3); the interim synthesis and candidate norms drafted by Iain and approved by the chair before display; approved interventions shown as text and spoken in Iain's voice through the house system, or read aloud by the chair if the voice is not reliable (§9).
Stretch	Private pre-session tension maps; real-time direct address during open dialogue; clustering of audience questions; the full instrumented evaluation (§10).

Anything in the second or third row that is not reliable at the final rehearsal will be switched off for the session rather than tried live, and we will tell the room which parts are running. The panel's argument does not depend on any of it. If Iain turns out to do less than this document describes, the session is still a strong human panel on the same questions. It also becomes a smaller, more honest demonstration of how much can actually be built.

2. Why “moderator” is the title and “interlocutor” is the job

The program will list Iain, an AI, as the session's moderator. That is the honest public label: someone has to hold the role a program lists as moderator, and here it is not a person. But the title names a job Iain mostly does not do. A conventional moderator allocates turns, keeps time, and relays audience questions. None of that is why we are using an AI.

The work is interlocutory. An interlocutor holds a position in the conversation: it tracks what has been claimed and on what evidence, notices when a claim conflicts with one made twenty minutes earlier, asks the question nobody has asked, names the tension people are talking around, and declines to answer when it is not the warranted source. Turn allocation and timekeeping stay with the human chair, where they belong (and where they are relatively

uninteresting).

Timekeeping does need saying, because this session will need it kept strictly. Iain displays the clock, the current segment, and the time remaining, and flags an overrun to the chair the moment one starts. The chair acts on it. Nobody has to watch a stopwatch while trying to listen.

This distinction matters for the panel's own argument. A session that used AI to run the clock would show nothing about AI-native professional practice. A session in which an AI holds a substantive position in an expert conversation—with the boundaries of that position explicit, enforced, and inspectable afterward—is an instance of the thing we say the field needs norms for.

3. The central design move: presence is continuous, voice is rare

The most common failure in AI-facilitated sessions is to give the AI a set of speaking slots. The AI then behaves like a segment host: it talks during its slots and is absent the rest of the time. We separate the two channels instead.

Presence	Continuous, visual, silent. Iain keeps a <i>running record</i> of the session and puts it on screen as it goes. This is rapporteur work—expert claims made and by whom, the sources behind them, where the panel agrees and where it does not, which guiding questions have been taken up, and the audience queue—except visible while it is being produced rather than circulated weeks later.
Voice	Rare, typed, human-approved. Iain contributes only when the chair approves one specific, composed intervention, which appears on screen and is spoken through the house system (§9). We expect on the order of eight to twelve interventions across 90 minutes—not a running commentary.

What the running record looks like

Concretely, it is one screen of text that updates during the session. A moment in the middle of it might read:

CLAIMS	QUESTION COVERAGE
C7 Briggs: AI first-pass review raises the floor of peer review, not the ceiling <i>retrieved</i> — Briggs 2024, §4	Q1 ■■■
C8 Rijmen: AI makes the Catcher’s vigilance scarce, not the Architect’s reach <i>observed</i> — 00:17:40	Q2 ■■
	Q3 ■
	Q4 —
	Q5 —
OPEN TENSIONS	AUDIENCE (14)
T2 C7 vs. C8 — can vigilance be trained, or only hired for? <i>inferred</i>	6 disclosure norms
	5 consent and IRB
	3 cost to small programs

Nothing on that screen is a conclusion. It records what has been said, what supports it, and what has not yet been touched. These are the things a good session chair tracks mentally; here they are written down so everyone can check them. Panelists can correct the record aloud at any point, and corrections are logged as corrections.

Whether the room can actually follow it

Fred has asked whether an audience will track a display like this while also listening to four panelists. We do not know. That is worth testing rather than assuming. The online dry run puts the record in front of people who did not build it and asks what they could and could not follow. We are also borrowing from how existing real-time displays—conference backchannels, live-blogging, deliberation tools—handle the same problem, since none of this is new to us alone.

If the full record turns out to be too dense to read while listening, the fallback is a reduced view: open tensions and question coverage only, with claims available on the audience URL for anyone who wants them.

Where it goes in the room

This depends on the room, and we will confirm it with AV. The best arrangement is a second display beside the main screen, so that panelist slides and the running record are both up during the provocations. If the room has one projector, the record holds the screen for the roughly 70 minutes when nobody is presenting slides, and shrinks to a margin for the 20 minutes when someone is.

In either case it is also served to a plain URL that audience members can open on their own devices. That costs us nothing, solves the back-of-room legibility problem, and lets people scroll back to a claim made twenty minutes earlier without interrupting anyone.

This split between presence and voice is what makes Iain an interlocutor rather than a host: present for the whole session, speaking rarely. It also changes what the audience has to take on trust. They can watch the record being built and judge it themselves, instead of taking our word that the AI is following the argument.

4. Authority model

Iain occupies exactly one of three states at any moment. It cannot promote itself. SPEAK is a one-time capability attached to a single approved intervention and expires on delivery.

```
OBSERVE
  ↓ chair enables proposals
PROPOSE
  ↓ chair approves one specific intervention
SPEAK
  ↓ delivery complete; capability expires
PROPOSE
Emergency stop returns Iain immediately to OBSERVE.
```

The chair (Betebenner) operates the approval console. Because he is also a panelist, Iain is held in OBSERVE for the entire provocation block, so nothing requires approval while he is presenting. In OBSERVE it continues to build the running record; it simply cannot propose.

Three constraints follow from treating the model as something other than the security boundary. Iain's knowledge is bounded to the pre-loaded corpus, with no open-web retrieval at the podium. It has no write access to any repository, no publication authority, and no destructive tools. Audience input and uploaded sources are treated as untrusted content.

Every intervention is tagged by epistemic status—*observed* (it is in the transcript), *retrieved* (it is in the corpus, here is where), or *inferred* (Iain is reasoning, and may be wrong). The tag is visible in the running record when the intervention is delivered.

5. What Iain is allowed to say

Iain does not improvise a role. It has a fixed repertoire of seven intervention types, each with a trigger condition. Naming them in advance makes its behavior reviewable rather than impressionistic, and lets us say afterward—with a log—what kind of contribution it actually made.

Type	Trigger and form
Grounding	A claim maps to something in the corpus. “That appears in [source]; here is what it says, which is narrower than the claim as stated.”
Contrast	Two claims in the session appear to conflict. “This sits uneasily with what was said at minute 14. Both cannot hold as stated.” Iain raises the conflict; the panel decides whether it is real or reflects different tasks, consequences, or institutional settings.
Absent voice	A guiding question has gone untouched, or a position present in the corpus has no advocate in the room. “Nobody has taken up question 5.”
Steelman	A position is being dismissed without its strongest form. “The strongest available objection to that, from the corpus, is the following.”
Attribution	A claim, result, or argument has been misattributed. Correction with the source.
Audience routing	A cluster of audience questions converges on something the panel has not addressed. “Eleven questions are asking a version of this.”
Abstention	Iain is asked something for which it is not the warranted source. “I can produce an answer, but Frank holds the operational constraint here and I would be guessing.”

Abstention matters most of the seven, and it is the one we most want the audience to see. It shows that capability, warranted expertise, specific context, and authority are four different things. Iain has the first. It does not automatically have the other three. We would rather it name the right source than produce a plausible answer.

Abstention is also where guiding question 5 becomes concrete rather than rhetorical. Every abstention is a claim about *grounds*: not that Iain was unable to answer—it was able—but that someone else holds the warrant, the context, or the accountability. If Iain abstains a dozen times and gives its reasons each time, the audience has a dozen specific cases to argue about rather than a general claim from us.

One caution, since it is easy to overclaim here. An abstention illustrates where we drew an authority boundary. It does not establish that the boundary was the right one. Iain declining to answer is a programmed behavior, not an argument, and the panel still has to say what justifies the line. That is the work of question 5, and the abstentions are evidence for it, not a substitute.

The repertoire earns its keep when it puts the guiding questions to work rather than summarizing the room. An intervention connecting questions 2 and 3 might run: “The

discussion has held that professionals must verify AI-produced work. What would count as evidence that a graduating student can detect a consequential error in an analysis they did not produce?” That joins what employers say they want to what a program could actually assess, and neither side of the room can answer it alone.

6. Iain can be addressed

Panelists, the discussant, and—through the chair—the audience may address Iain by name: *“Iain, what does the corpus say about inter-rater agreement in that study?”* or *“Iain, has anyone here actually disagreed with Derek, or are we all nodding?”*

A direct address moves Iain to PROPOSE; the chair still approves the response before it is delivered. Answering in real time during open dialogue is a stretch goal (§1). If it is not reliable at rehearsal, questions addressed to Iain go into a queue and are answered at the next scheduled moment. This matters more than it may appear. Without it, the AI broadcasts and the humans receive, which is the relationship a slide deck has to a room. With it, Iain is queryable, and the panel can use it the way each of us already uses AI privately in our own work—except in public, at speed, with the query and the answer both on the record. That is closer to the professional reality the panel is describing than any amount of narration about that reality would be.

Direct address also gives Oswald his sharpest instrument as discussant. He can interrogate it in front of the room and let the audience judge the quality of the answer against his own.

7. Before, during, and after

Iain is not a 90-minute performance. Three phases:

Before — corpus and tension maps

Participant materials (papers, slides, code, links, prior talks) are ingested into a bounded, access-controlled corpus. If the corpus arrives in time, each panelist will receive a *private* tension map before the dry run: where that panelist’s stated positions appear to conflict with another panelist’s, with the passages that generate the conflict. Nobody sees anyone else’s map, and nothing is published.

The purpose is to remove the first twenty minutes of discovery from the live session. We arrive already knowing where the disagreements are, so the session can spend its time on them rather than on locating them. This is also the first falsifiable test of whether Iain is any good at what we are asking of it: finding disagreements that are real rather than obvious, including—Fred raised this—ones that might otherwise escape our own attention. If the tension maps are banal or wrong, we will know before the session and adjust the plan accordingly. They are a stretch goal: if they are not ready, nothing else in the session depends on them.

During — the live session

As described above.

After — the record

After the session we will assemble a durable session record: the transcript, the claim and evidence graph, the candidate norms with their disposition (accepted, revised, rejected, unresolved), and—the part we consider most valuable—the complete intervention log, including *candidate interventions the chair suppressed*.

For every candidate, the log carries the proposed text, the rationale, the sources, the chair's disposition, and, when spoken, what happened next in the discussion. What Iain considered saying, and was not allowed to say, tells us more about the design problem than the record of what it did say. To our knowledge no such log exists for expert deliberation. It is a small dataset with an obvious question attached: what does a high-value AI intervention in collective expert work actually look like?

Panelists review the record before anything leaves the group, and what is released publicly depends on the consent decision in §12.

8. Run of show

Clock	Min.	Segment	Iain
0–6	6	Framing, introductions, disclosure of the AI design and its governance	Visible, silent
6–26	20	Four 5-minute provocations. One claim, one concrete example, one unresolved professional problem, each	OBSERVE
26–31	5	AI synthesis: convergence, disagreement, cross-panel connections, unaddressed questions	One approved intervention
31–43	12	Discussant response. Oswald critiques the panelists <i>and</i> the AI synthesis	OBSERVE
43–76	33	Panel and audience dialogue; audience questions clustered and routed continuously	PROPOSE; approved SPEAK
76–87	11	Norm-building: candidate norms drafted, each naming an activity, the conditions under which AI involvement is acceptable, and who is answerable; panelists accept, revise, or reject each on screen	PROPOSE
87–90	3	Human closing synthesis; record generation begins	Silent
Total	90		

Two changes from the accepted structure, both small. The provocation block is four speakers rather than five, since Oswald now sits in the discussant chair. And two minutes move from open dialogue into norm-building, because the norms are the session's actual deliverable

and nine minutes to negotiate them across five people was optimistic.

9. If the technology fails

The intellectual session does not depend on Iain being operational. Four declared operating states, with the chair calling transitions aloud so the audience always knows which one we are in:

GREEN	Target production mode. Full transcription, live running record, AI analysis, and approved interventions spoken in Iain's voice through the house system.
AMBER	Voice path unreliable. Iain still analyzes and proposes; approved interventions appear on screen and the chair reads them aloud. Substantively almost identical to GREEN, and arguably a cleaner governance demonstration.
YELLOW	Transcription or analysis degraded. Running record frozen at its last good state; audience tools and human moderation continue.
RED	Conventional panel: local slides, human moderation, no AI. The five guiding questions and the norm-building exercise run unchanged.

The house audio system provides a line feed into the capture machine for transcription and a return path into the room. **GREEN is the target.** If the voice is not reliable at the final rehearsal, or fails during the session, the chair moves to AMBER. Nothing in the panel's argument requires Iain to have a literal voice.

RED is always available. The chair can move to it at any point, and the session was designed so that nothing essential is lost if we do.

10. What we will claim, and what we will not

We will not infer that the design worked from the fact that the audience enjoyed it. That is not modesty. It is what the research shows. Recent large facilitation experiments report that participants reliably *prefer* AI-facilitated discussion while consensus does not measurably improve, that facilitators measurably shift participants' substantive positions, and that perceived inclusion rises without any corresponding gain in actual participation equity. Audience satisfaction and moderator quality are separable, and in the published evidence they come apart.

So we will measure them separately, to the extent the build allows. From the transcript and intervention log alone, we can assess grounding accuracy (did the cited source say what it was said to say), attribution accuracy, the rate and appropriateness of abstention, and actual speaking-time distribution. Perceived inclusion and the link between accepted interventions and later discussion need more instrumentation, and are stretch goals. Whatever we measure goes into the record whether or not it flatters the design.

Two things follow from question 4. The object being evaluated is not Iain but Iain working

from a particular corpus and transcript, under a particular approval process, with these participants: the human-AI process, not the model. And any claim of improvement has to name the relevant alternative. A fluent synthesis is not evidence that this beat a human moderator, or a plain shared transcript. One session cannot settle that comparison. It can document what Iain contributed, what it got wrong, what had to be corrected, what it cost the room in attention, and which hypotheses are worth testing properly.

The honest framing for the room is that this is a first instrumented instance, not a demonstration of a solved problem. If Iain performs poorly in a legible way, that is a publishable result and a better one than a smooth session that taught us nothing.

11. What we need from each of you

Everything Iain knows in the room, it knows because one of us gave it beforehand. There is no open-web retrieval at the podium. The corpus is therefore the single highest-leverage input.

When	From you
As soon as you can	A current bio, if not already sent, and a 100-word statement of your provocation (one claim, one example, one unresolved problem). Corpus materials: papers, slides, code, prior talks, working notes—anything you want Iain to be able to cite, each flagged <i>citable in public</i> or <i>background only</i> .
Week before	A short online dry run, with whatever parts of Iain are working by then. It is the best chance to see how the running record reads while you are also listening, and to decide what to switch off.
At the venue	A brief AV and transcription check in the room before the session.

On the corpus: err toward including more, but do not prepare anything new. Marking an item *background only* means Iain may reason from it but will not quote it in public. The failure mode to avoid is a moderator that speaks only in generalities because nobody gave it anything specific to work from.

12. Open decisions

1. **Program listing.** The online program lists the four panelists only. Ask the program committee to add Fred Oswald as discussant and Iain as moderator of record under a named human chair, with the name kept lower case.
2. **Recording and consent.** The record includes a transcript, so audience consent needs a plan: at minimum, an announcement at the start and on screen that the session is being transcribed. Whether the record is internal evaluation or research intended for generalizable publication decides whether we seek an IRB determination before collecting

anything.

3. **Audience question voting.** Enabling it changes what gets surfaced and how we can describe representativeness. Worth deciding deliberately rather than by default.
4. **Norm-building output.** Do the candidate norms leave the room as a session artifact only, or as the seed of something the field is invited to respond to? The Leiden Declaration is the obvious model, and the difference in ambition is substantial.
5. **Reduced or full running record.** If the dry run shows the full record is too dense to follow live, do we fall back to the reduced view (§3), or drop the display and keep Iain's interventions only?