

Educational Measurement as an AI-Native Profession

AIME-CON 2026 · PANEL · WEDNESDAY 7 OCTOBER, 11:00–12:30 · COMMONWEALTH 2

Session chair Damian Betebenner, Center for Assessment

Moderator Iain, an AI interlocutor under the chair's supervision

Discussant Fred Oswald, University of California, Irvine

Panelists

Derek Briggs, University of Colorado Boulder


Frank Rijmen, Cambium Assessment

Mohammed A. A. Abulela, MetaMetrics / University of Minnesota

Damian Betebenner, Center for Assessment

A correction to the printed program: it lists a human moderator. Today the moderator is Iain, an AI, under the session chair, and Fred Oswald is our discussant.

What this conference is about



Of the 335 items in the printed program, **66%** build AI into what we deliver and **39%** already use AI in how we work. **22 items (7%)** ask what that means for the profession itself.

Rows overlap: each item was coded on three independent yes/no questions by an AI coder, and a blind second AI coder on a stratified sample of 98 agreed on 91%, 97% and 94% (κ 0.80, 0.94, 0.72). Audited by the chair. Data, prompts and agreement: the materials site.

The provocation

At this conference: products **221** · our methods **131** · the profession itself **22** (of 335 items).

THE CHAIR'S CONJECTURE: AN INDEPENDENT ORDER-OF-MAGNITUDE CLAIM, NOT ESTIMATED FROM THE PROGRAM CODING

How measurement professionals use AI in their own work will have 10× the impact of the AI we build into assessments.

AI in the product changes one tool: a scoring engine, an item pool, a tutor.

AI in the profession changes every analysis, model, review, document, and decision, including how the next generation of those products is designed, validated, and governed.

Why a profession-level response

The mathematics community has begun asking this explicitly. The [Leiden Declaration on AI and Mathematics](#) (June 2026, endorsed by the IMU; [leidendeclaration.ai](#)) asks the discipline to protect the verifiability of proof, attribution, authorship, and the primacy of human judgment.

Five guiding questions

1. **What has changed** in the substance and organization of measurement work?
2. **What that demands of professionals:** the division of labor, expertise, and training.
3. **What the profession must redefine:** what is verified, documented, and disclosed.
4. **What measurement science owes the evaluation of AI**, including this panel.
5. **Where authority and accountability remain human**, and on what grounds.

Meet Iain

AN EXPERIMENT IN AI-NATIVE EDUCATIONAL MEASUREMENT

Iain is our attempt to show what AI-native educational measurement could look like: in this room, not on a slide.

Its evidence is bounded to the panelists' materials and what is said in this room, with no open web. Like any language model it carries pretrained knowledge, but it may not present that as evidence: a source claim needs a citation, a claim from the room needs a timestamp, and anything else is labeled inferred. Its authority belongs to the chair. Everything it proposes, including what it is not allowed to say, goes on the record.

As AI becomes more capable, one of the most important things professionals do is this: gather, deliberate in public, and decide what counts as warranted.

A session like this one can leave more than memories. The materials, the claims, the disagreements and the norms are collected into a record that future colleagues, and the AI systems they work with, can query, for example through MCP or other AI connectors.

How Iain takes part

Presence: continuous, visual, silent

A running record on screen: claims and who made them, sources, agreements, open tensions, question coverage.

Voice: rare, typed, human-approved

About 8–12 interventions in 90 minutes, each approved by the chair.

Every contribution is tagged **observed** (the transcript records someone saying it, with a time), **retrieved** (in a panelist's material, with the source), or **inferred** (Iain's own reasoning, possibly wrong). Its evidence: the panelists' materials and the transcript; no open web. It **abstains** when someone else holds the warrant.

OBSERVE

↓ chair enables proposals

PROPOSE

↓ chair approves one intervention

SPEAK

↓ delivered; the permission expires

Emergency stop → OBSERVE, at once

If the technology fails

State	What it means
GREEN	Live transcription and running record; approved interventions spoken in Iain's voice
AMBER	Same, but I read approved interventions aloud
YELLOW	Transcription or analysis degraded; the record freezes at its last good state
RED	A conventional panel: same questions, same norm-building

Running today: I'll tell you which state we're in, and I will call every change of state aloud. The panel's argument does not depend on Iain working.

Before we start

This session is being recorded and transcribed.

lain, an AI moderator, speaks aloud under the session chair, who approves every sentence. The transcript and the session record are published afterwards. Questions, typed or asked aloud, become part of that record. No names or emails are collected.

The record includes the full intervention log, including what lain proposed and I declined. Panelists review it before release.



Follow along · ask a question
or open it from the materials site



Panel materials
dbetebenner.github.io/AIME_2026_Panel_Presentation

Run of show

Minutes	Segment	lain
0–6	Framing and the AI design	visible, silent
6–26	Four five-minute provocations: one claim, one example, one unresolved problem	OBSERVE
26–31	AI synthesis: convergence, disagreement, connections	one approved intervention
31–43	Discussant: Fred Oswald critiques the panel and the AI synthesis	OBSERVE
43–76	Panel and audience dialogue	PROPOSE; approved SPEAK
76–87	Norm-building: activity · conditions for AI involvement · who is answerable	PROPOSE
87–90	Human closing synthesis	silent